



# What is Docker





# Docker Internals

- Container are method of operating system virtualization that allow you to run application and its dependencies in resource isolated process.
- Containers allows you to easily package an application code, configurations and dependencies into easy to use building blocks and that application can be deployed quickly and consistently regardless of deployment environment.
- Linux kernel features that create the walls between container and other processes running on the host.

To understand containers, we have to start with Linux Cgroup and Namespace.

**Namespace** - Wrap a set of system resources and present them to a process to make it look like they are dedicated to that process.

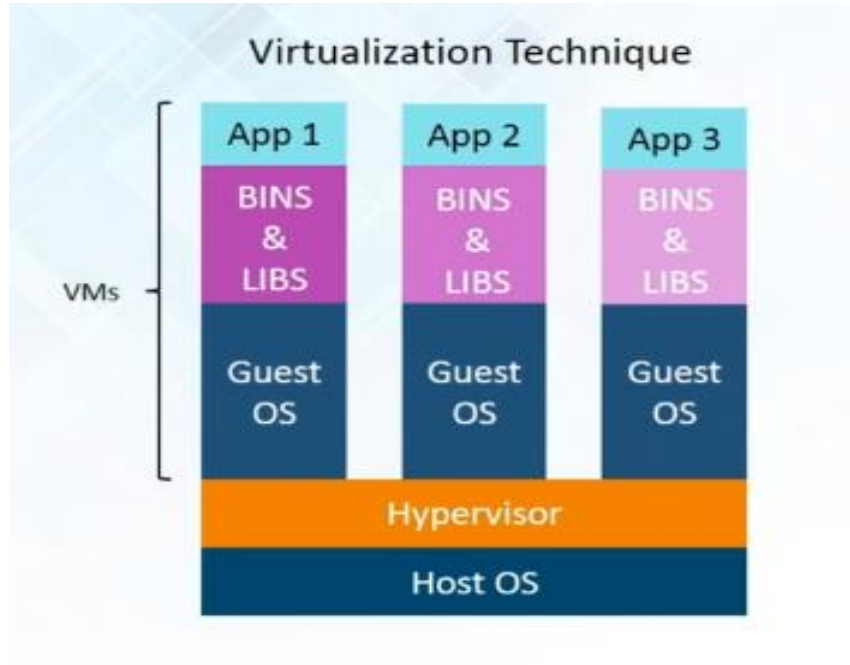
## Docker Internals

**Cgroup** - Governs the isolation and usage of system resources such as cpu and memory for group of process. E.g. If you have a application that takes up lot of cpu cycles and memory such as scientific computing application you can put the application in a cgroup to limit a CPU and memory usage.

**Namespaces** deal with resource isolation for single process

**Cgroup** manages resources for group of processes

# Virtualization




## Advantages

- Multiple OS In Same Machine
- Easy Maintenance & Recovery
- Lower Total Cost Of Ownership

## Disadvantages


- Multiple VMs Lead To Unstable Performance
- Hypervisors Are Not As Efficient As Host OS
- Long Boot-Up Process ( [Approx. 1 Minute](#) )

# Virtualization




# Virtualization

## One application on one physical server



- Slow deployments
- Huge costs
- Wasted resources
- Difficult to scale, migrate
- Vendor lock in

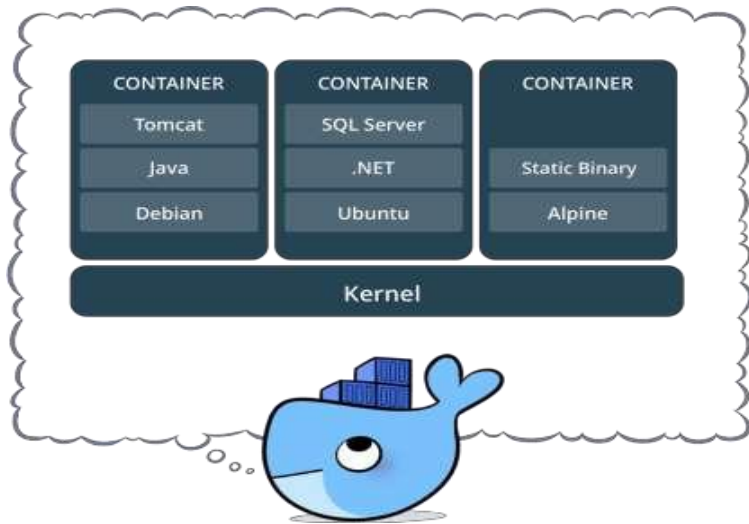
## Hypervisor Based Virtualization



- Better resource pooling  
One physical machine divided into multiple virtual machines
- Easier to scale
- VMs in the cloud  
Rapid elasticity  
Pay as you go model

# Why Docker

- Each VM still requires  
CPU allocation, Storage, RAM, Guest OS
- The more VMs you run, the more resources you need
- Application portability not guaranteed



- Standardized packaging for software and dependencies
- Isolate apps from each other
- Share the same OS kernel
- Works with all major Linux and Windows Server



# Transformation



Applications are transforming

2000

Today



# Applications Modernisation



## Developer Issues:

- Minor code changes require full re-compile and re-test
- Application becomes single point of failure
- Application is difficult to scale

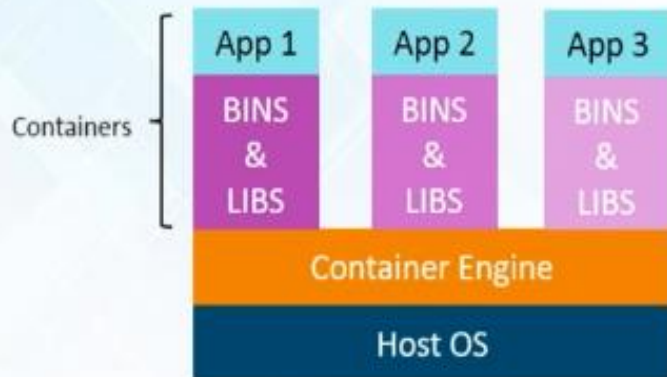
**Microservices:** Break application into separate operations

Make the app independently scalable, stateless, highly available by design

# Docker Info

Note: Containerization Is Just Virtualization At The OS Level

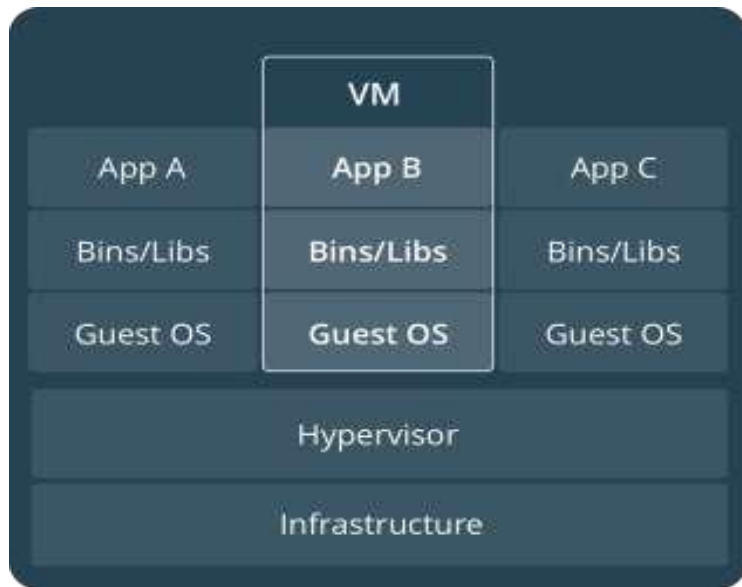
## Containerization Technique



## Advantages Over Virtualization

- Containers On Same OS Kernel Are Lighter & Smaller
- Better Resource Utilization Compared To VMs
- Short Boot-Up Process (  $\frac{1}{20^{\text{th}}}$  of a second )

## Containers vs VMs

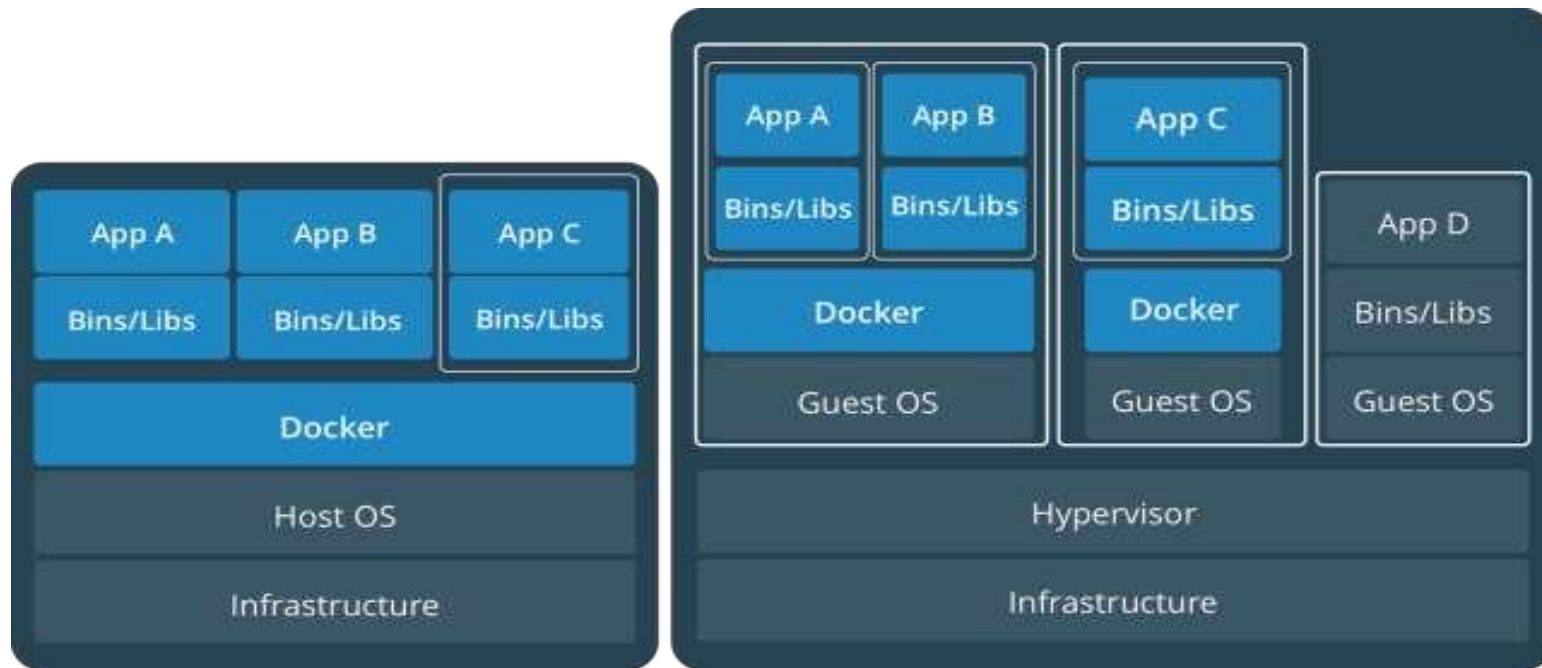


VMs are an infrastructure level construct to turn one machine into many servers



Containers are an app level construct



## Containers vs VMs



Containers and VMs together provide a tremendous amount of flexibility for IT to optimally deploy and manage apps.


# Docker Info

## In case of Virtual Machines



7 Gb of Memory is blocked and cannot be allotted to a new VM

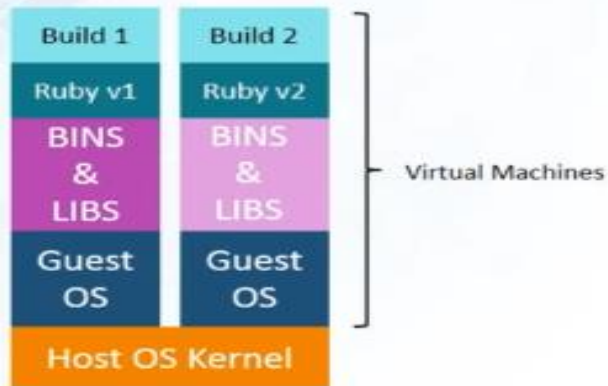
## In case of Docker



Only 9 GB memory utilized;  
7 GB can be allotted to a new Container


# Docker Info

## In case of Virtual Machines



New Builds → Multiple OS → Separate Libraries  
→ Heavy → More Time


## In case of Docker



New Builds → Same OS → Separate Libraries  
→ Lightweight → Less Time

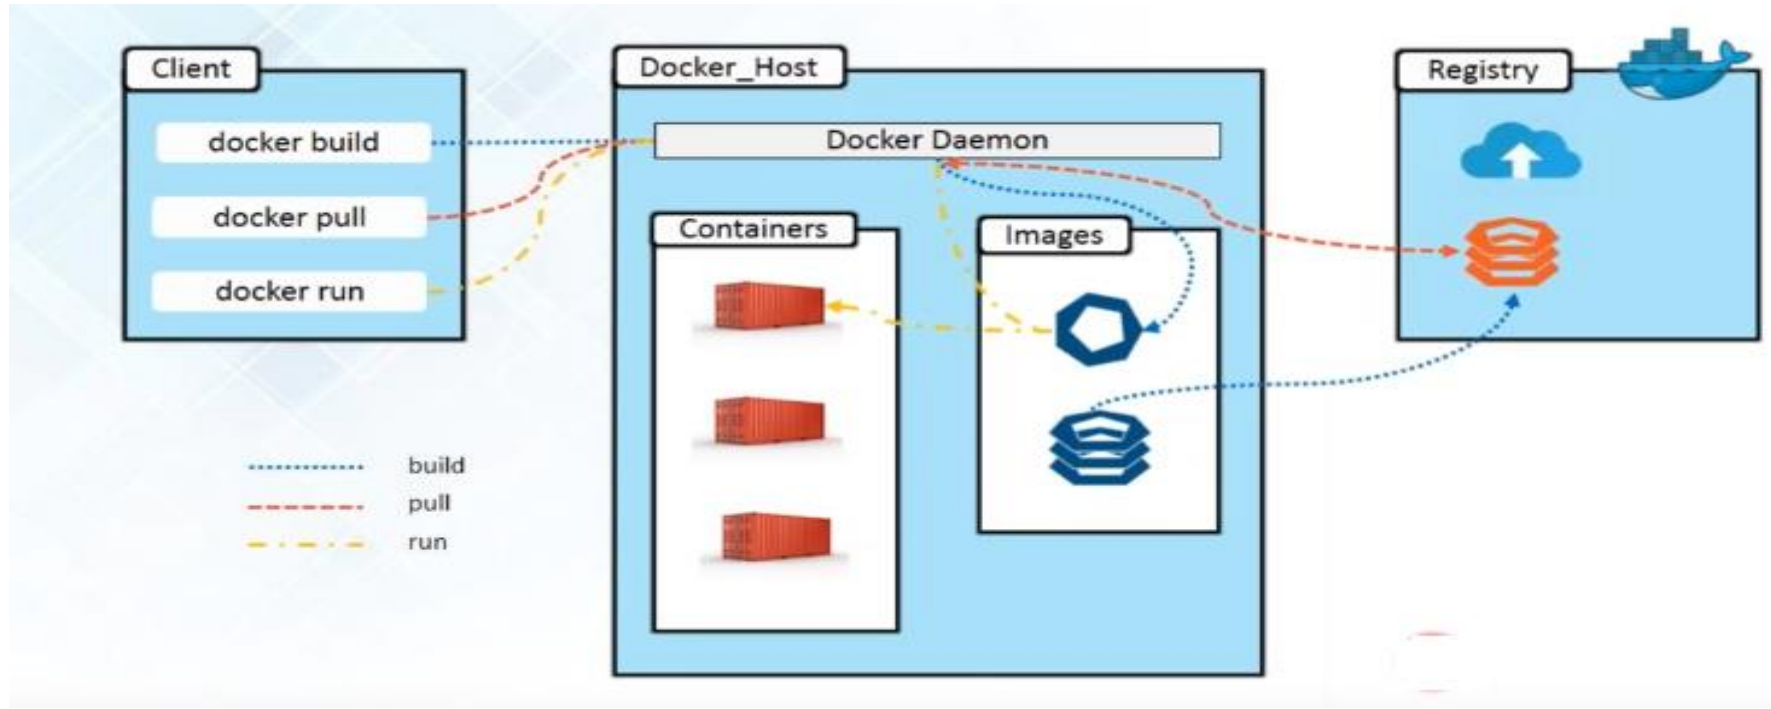
## Docker Information

Docker is a Containerization platform which packages your application and all its dependencies together in the form of Containers so as to ensure that your application works seamlessly in any environment be it Development or Test or Production.





# Docker Architecture



# Docker Registry

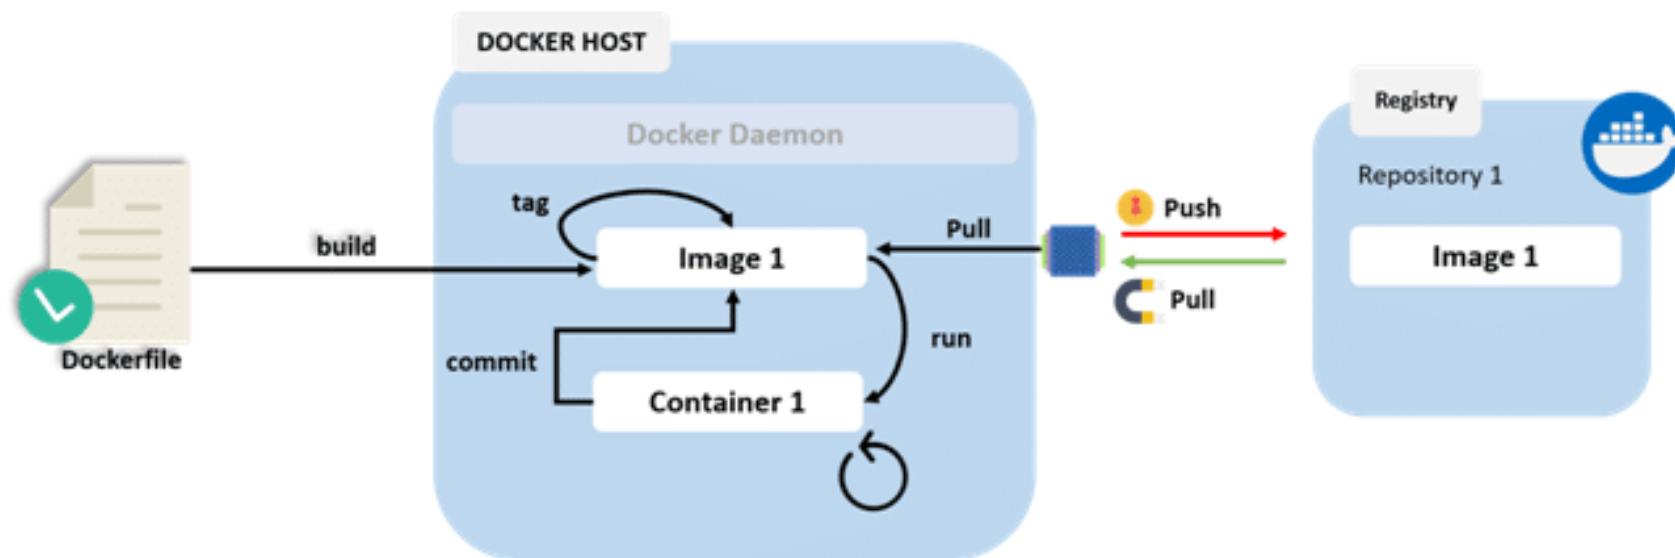
- Docker Registry is a storage component for Docker Images
- We can store the Images in either Public / Private repositories
- [Docker Hub](#) is Docker's very own cloud repository



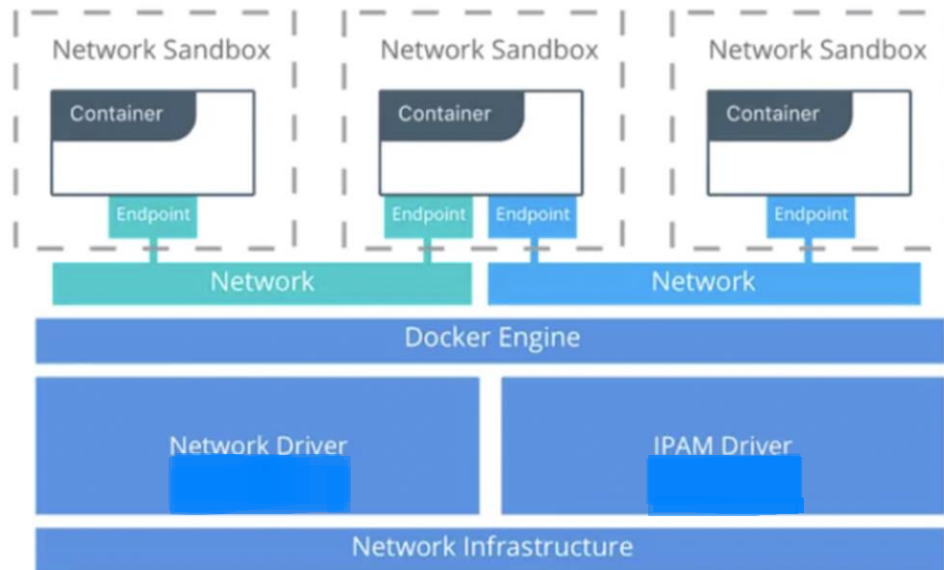
## Why Use Docker Registries?

- Control where your images are being stored
- Integrate image storage with your in-house development workflow

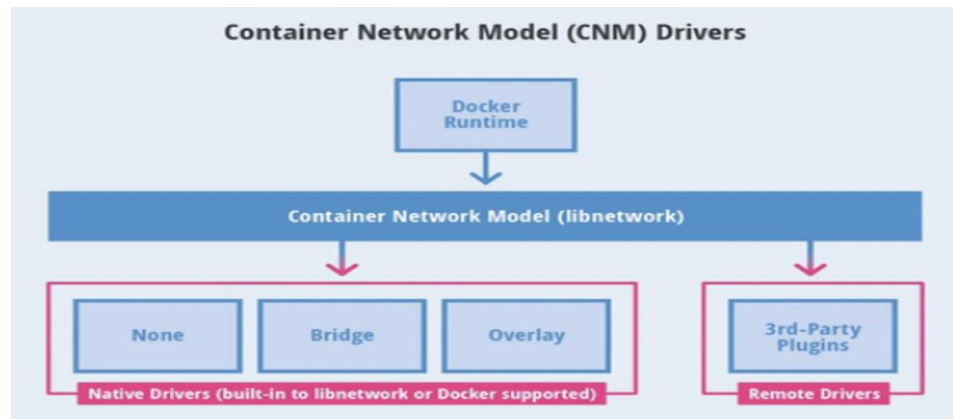
# Typical Workflow



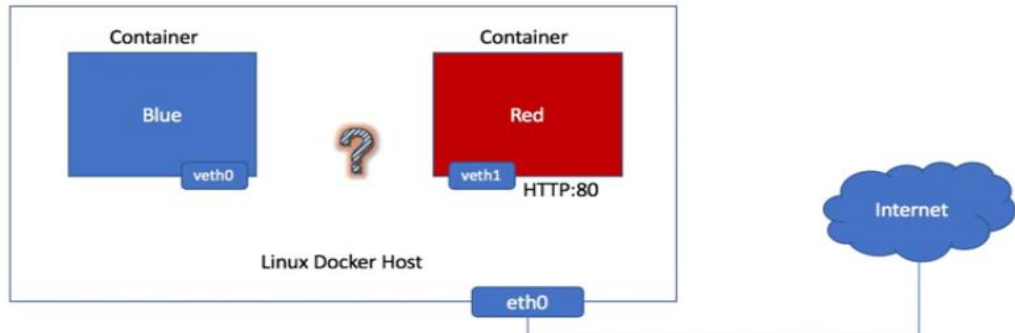
# The Container Network Model



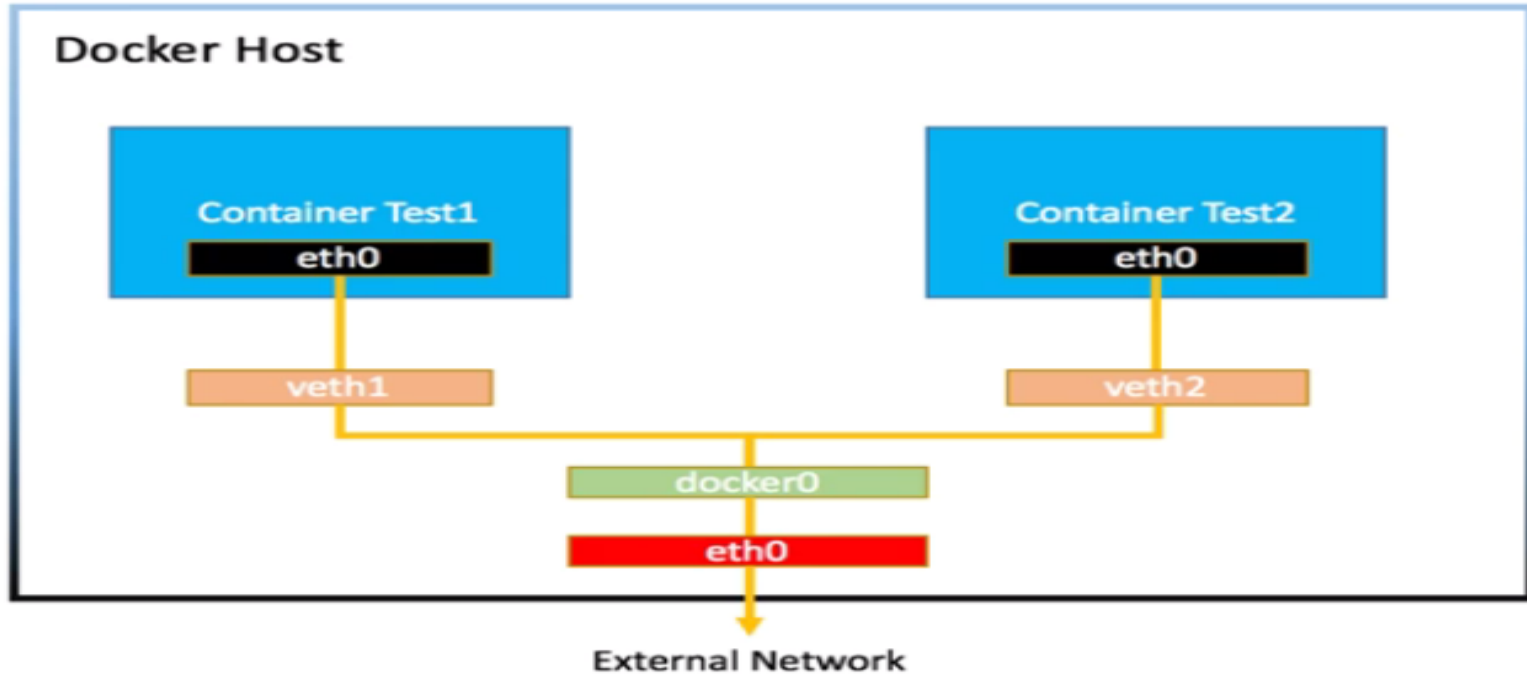
# Docker Info




- How two different containers in the same host communicate with each other?
- How the container communicate with the outside of Linux host (Internet)?
- How access container from the local docker host and outside of the docker host?



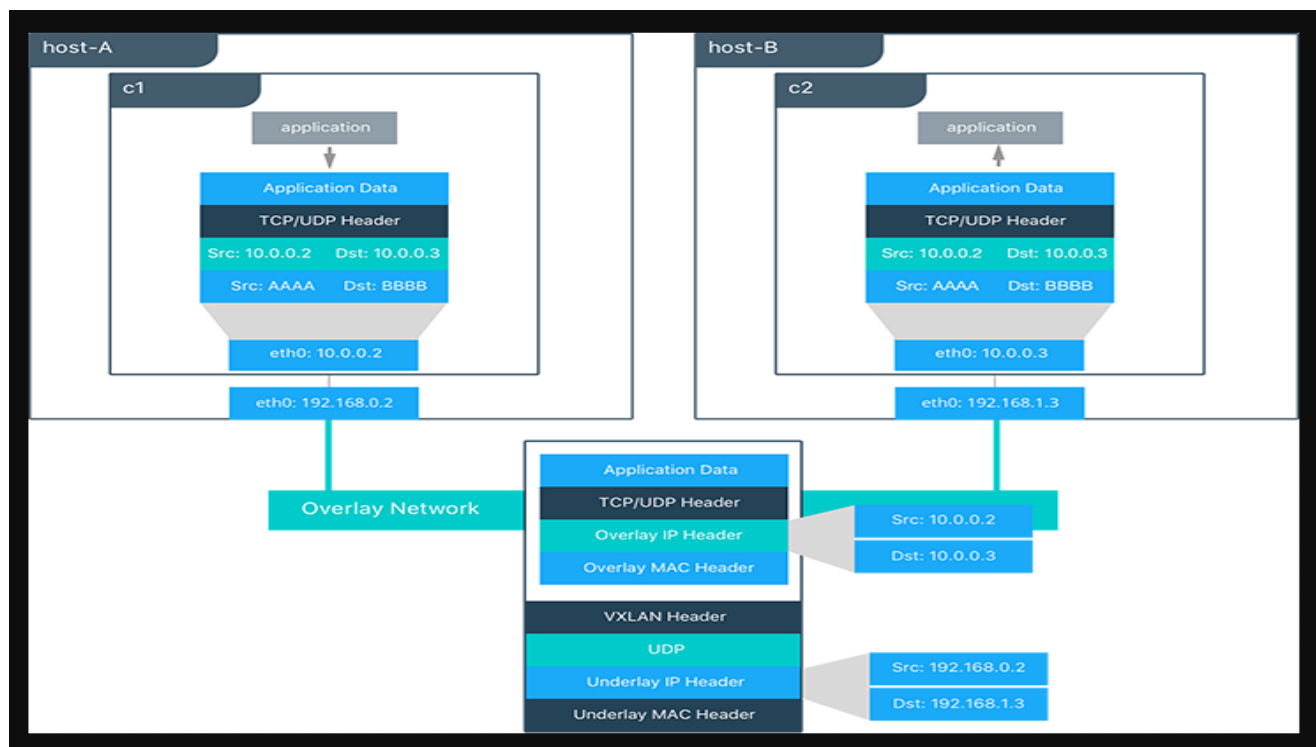
## Docker Info



# Multi host Networking



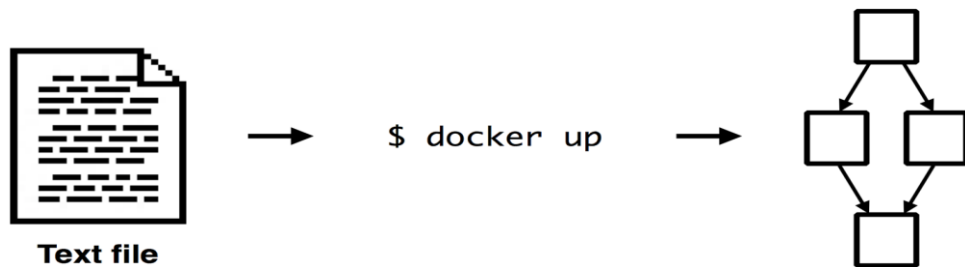
# Multi host Networking





# Docker Compose

- Multi-container apps are a hassle
- Build images from Dockerfiles
- Pull images from the Hub or a private registry Configure and create containers
- Start and stop containers



- Compose is a tool for defining and running multi-container Docker applications. With Compose, you use a Compose file to configure your application's services. Then, using a single command, you create and start all the services from your configuration

# app.py

```
from flask import Flask
from redis import Redis
import os
import socket

app = Flask(__name__)
redis = Redis(host=os.environ.get('REDIS_HOST', 'redis'), port=6379)

@app.route('/')
def hello():
    redis.incr('hits')
    return 'Hello Container World! I have been seen %s times and my hostname is %s.\n' %
(redis.get('hits'),socket.gethostname())

if __name__ == "__main__":
    app.run(host="0.0.0.0", port=5000, debug=True)
```

# Dockerfile

```
FROM python:2.7
MAINTAINER "abc@gmail.com"
COPY . /app
WORKDIR /app
RUN pip install -r requirements.txt
EXPOSE 5000
CMD [ "python", "app.py" ]
```

```
requirements.txt
flask
redis
```

- **FROM** : Every Dockerfile starts with this keyword. Put the image name after the keyword which to be a base image of yours. ex> FROM centos
- **ARG** : When **docker build** , able to get parameters. You can use something like **--build-arg <ARG name>=<value>**
- **ENV** : Environment variables for the docker image. i.e> JAVA\_HOME, CATALINA\_HOME
- **RUN** : Will executed when we `build`
- **EXPOSE** : port to be exposed to the host
- **VOLUME** : Define shared volumes
- **CMD** : Will executed when `**docker run**`


# Docker Info

```
version: "2"

services:
  web:
    build: .
    ports:
      - "80:5000"
    links:
      - redis
    networks:
      - compose-demo-bridge

  redis:
    image: redis
    ports: ["6379"]
    networks:
      - compose-demo-bridge

networks:
  compose-demo-bridge:
```



# Why we need Container Orchestration?

---

Do you need scaling beyond one host?

Do you need high availability?

Are your containers truly stateless?

## Production requirement

---

Availability – it just has to be working all the time, with as little downtimes as possible

Performance – our server needs to handle traffic, so performance is important

Easy deployment & rollback

Gathering logs & metrics

Scheduling: Where do these containers run?

Management: Who manages their life cycle?

Service Discovery: How do they find each other?

Load Balancing: How to route requests?

# Solutions

---

Kubernetes

Docker Swarm

Amazon EKS

AKS

GKE

# Kubernetes Overview

---


- Started by Google in 2014. First released in 2015
- Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications.
- Donated to Cloud Native Computing Foundation
- 100% open source
- Written in Go language



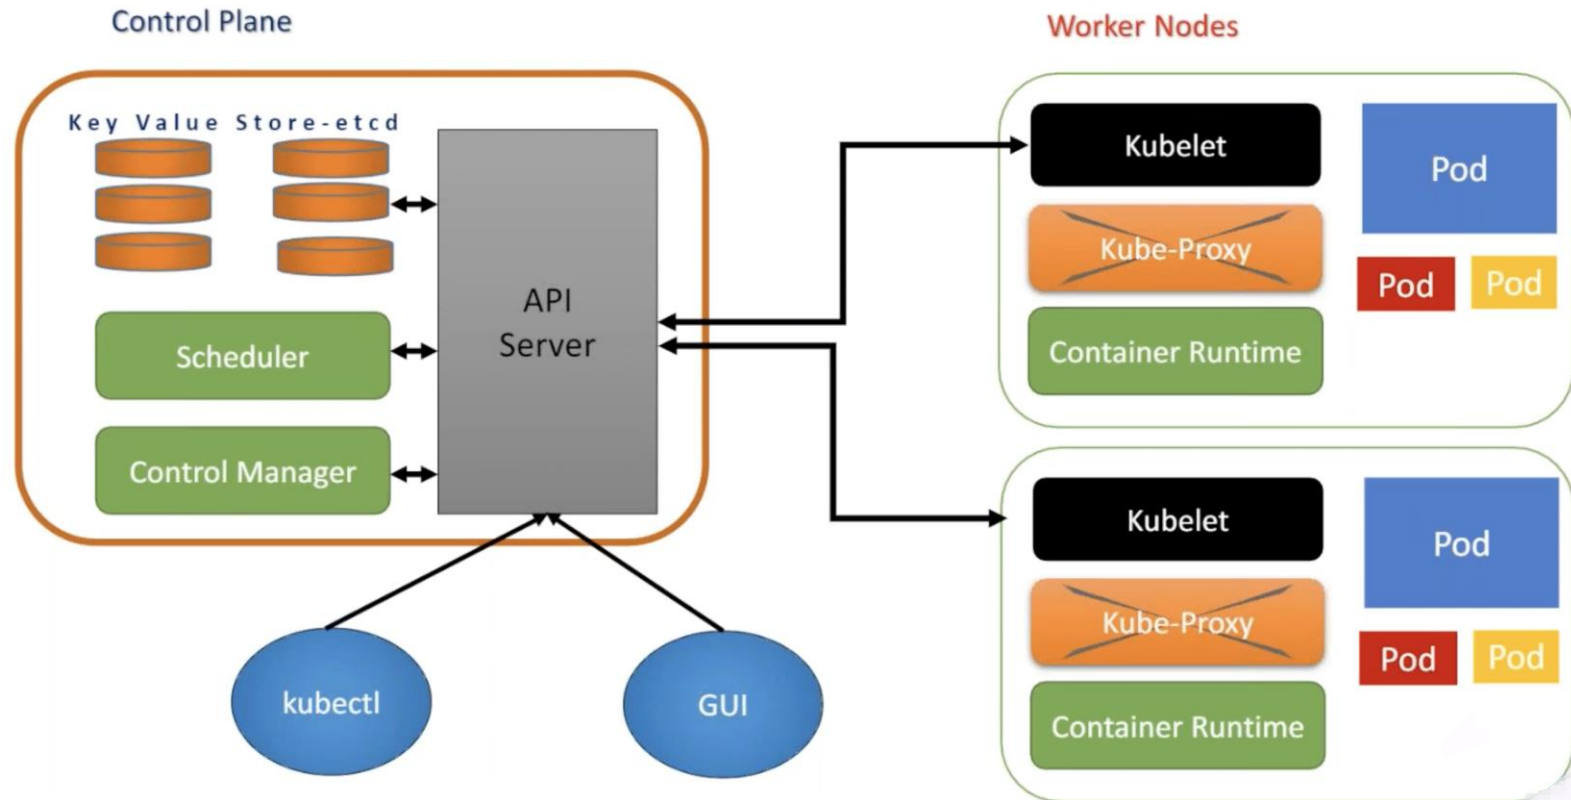
## Kubernetes – Basic Terms

- Kubernetes -- the whole orchestration system. K8s in short
- Kubectl -- CLI to configure Kubernetes and manage apps. Kube control
- Node -- Single server in the kubernetes cluster
- Kubelet -- Kubernetes agent running on nodes
- Control plane -- Set of containers that manage the cluster
  - Includes API server, scheduler, controller manager, etcd, etc
  - Also called master


# Kubernetes Architecture



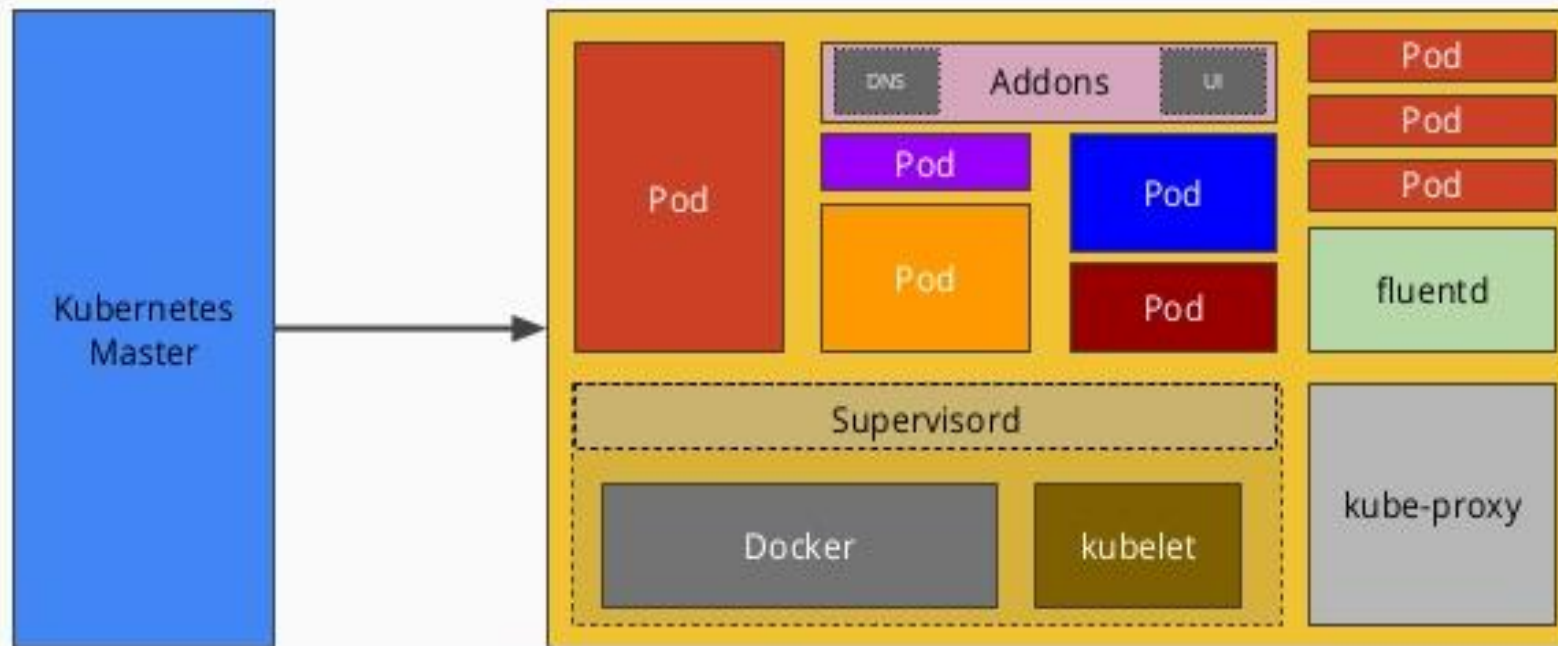
# Kubernetes Architecture



## Kubernetes Architecture - Master



## Kubernetes Architecture - Node



# POD

## Pod - Config

```
# nginx-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: nginx
    tier: dev
spec:
  containers:
  - name: nginx-container
    image: nginx
```

Kind	apiVersion
Pod	v1
ReplicationController	V1
Service	v1
ReplicaSet	apps/v1
Deployment	apps/v1
DaemonSet	apps/v1
Job	batch/v1

DEMO:-Pod Creation

# Kubernetes Basics

---

- **Nodes:** Hosts that run Kubernetes applications
- **Cluster:** Pool of compute, storage and network resources
- **Containers:** Units of packaging
- **Labels/Selector:** Key-Value pairs for identification
- **Pods:** Units of deployment
- **Deployment:** Auto Pod management
- **Replication Controller:** Ensures availability and scalability
- **ReplicaSet:** Advanced Replication controller with help of regular expression
- **Services:** Collection of pods exposed as an endpoint
- **Volumes:** To store data permanently
- **Secrets/ConfigMap:** To pass secrets and config files
- **DaemonSet:** To run similar task pod on each Node

# Kubernetes Services

- Exposing PODs to the outside world. Services are required for discovery.
- Labels and selectors are used to route to the appropriate POD
- Service Types
  - ClusterIP - Services makes internal pod accessible
  - NodePort - Service does default load balancing
  - LoadBalancer - Service does load balancing with external load balancer

**Ingress** - Path based load balancing with service type LoadBalancer and NodePort

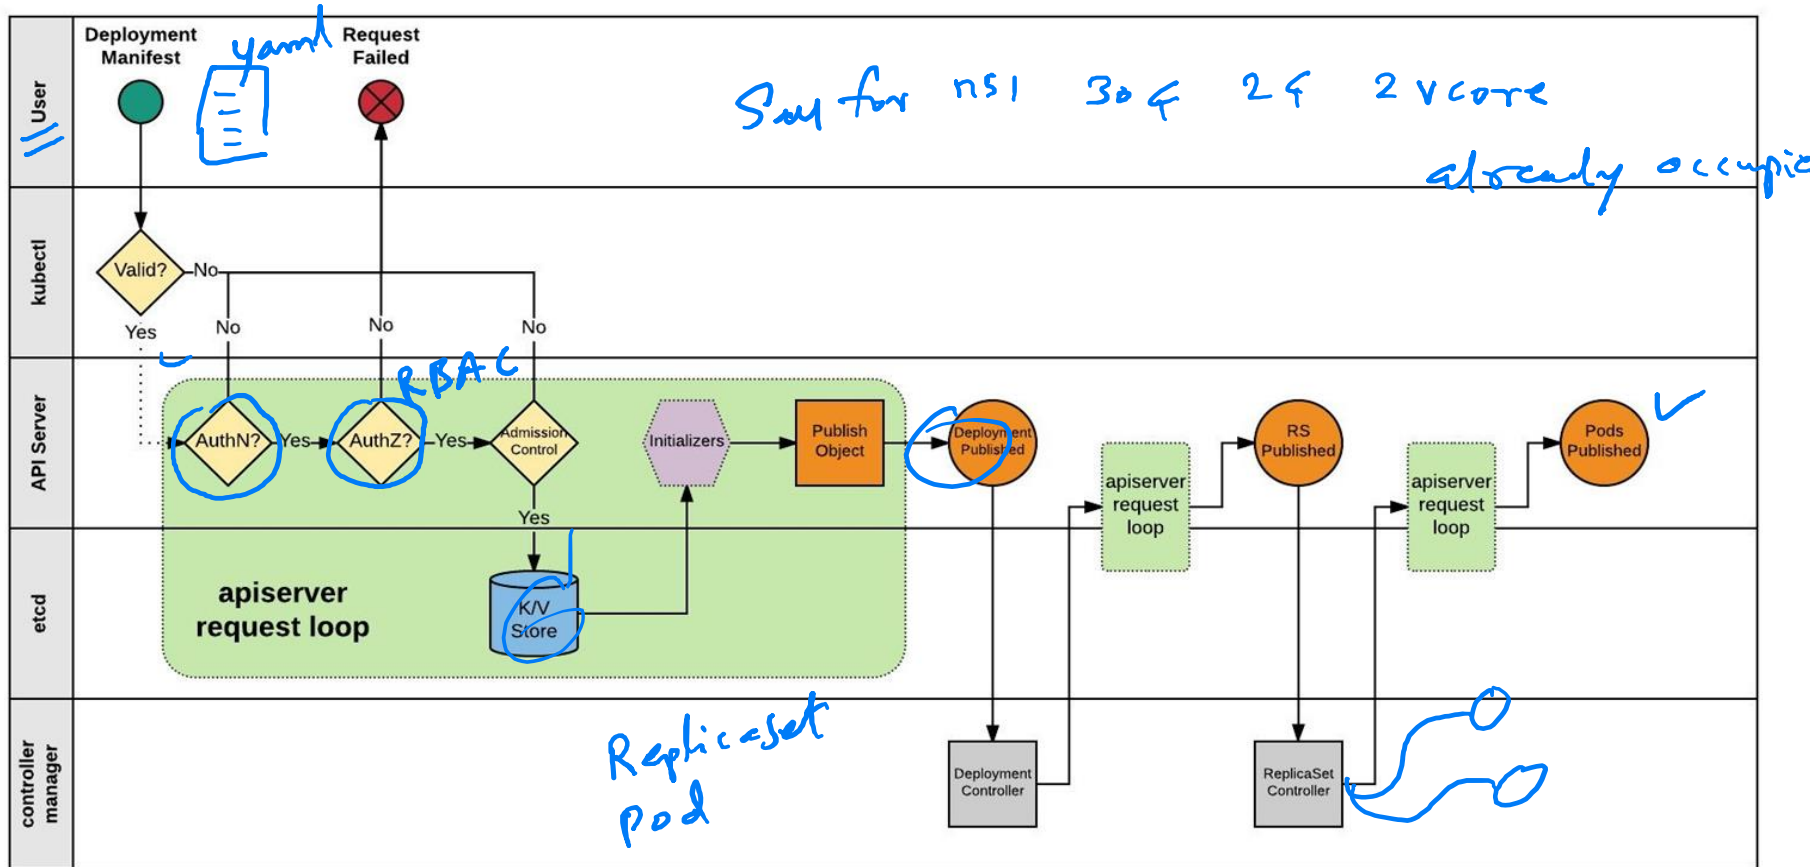


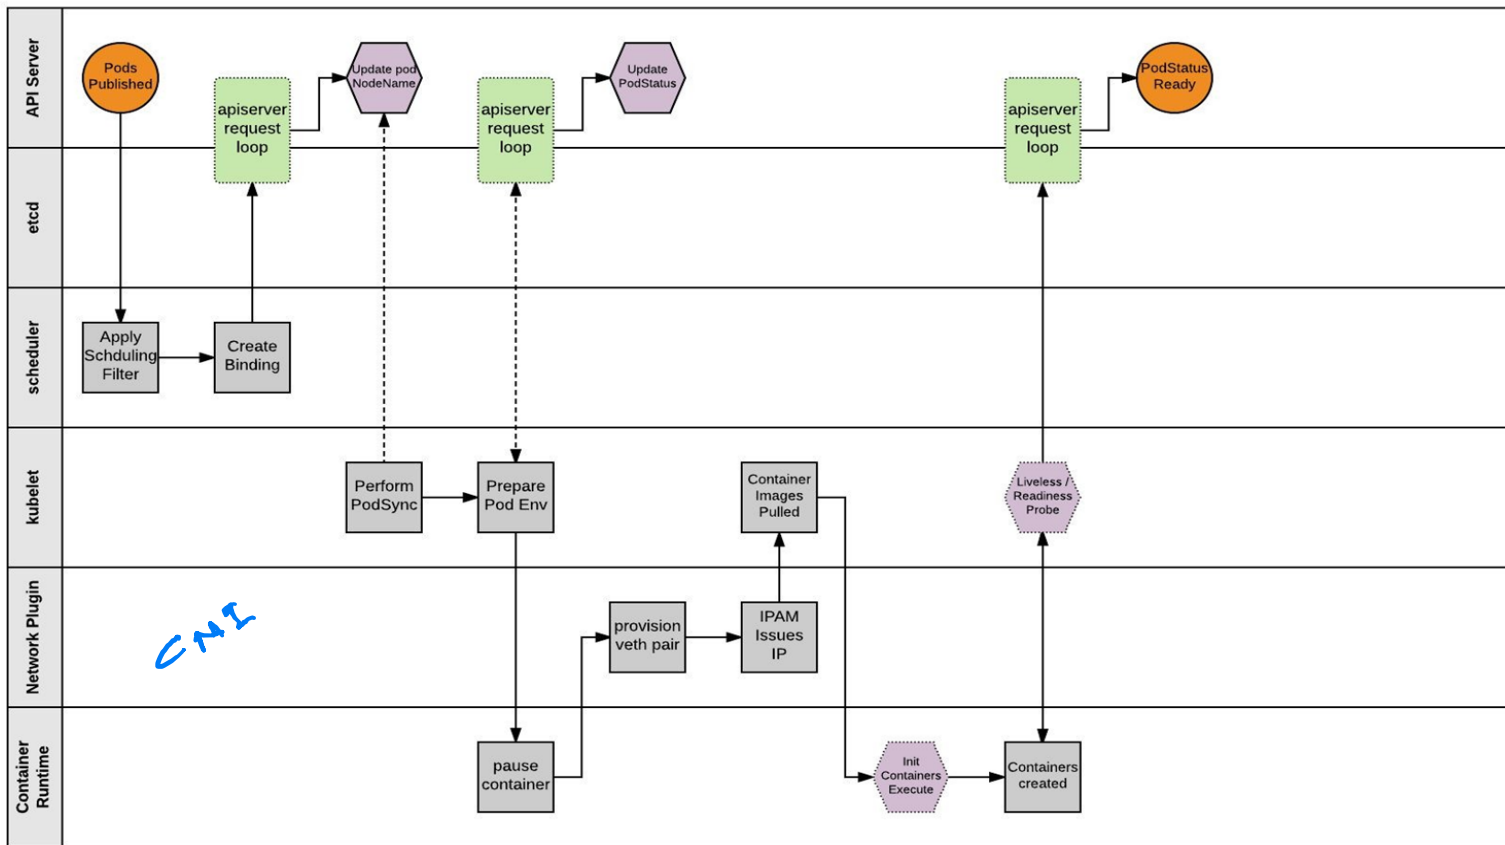
## Kubernetes - Helm Charts

---

- Makes deployment easy. POD management, configuration etc, separately managed than manifest
- Package manager for kubernetes
- Single command for installing, upgrading and deleting releases.

# Workflow





## Multinode cluster setup

Run below steps on master and slave machine

```
apt-get update && apt-get install -y apt-transport-https
```

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -
```

```
cat <<EOF > /etc/apt/sources.list.d/kubernetes.list
```

```
deb http://apt.kubernetes.io/ kubernetes-xenial main
```

```
EOF
```

```
apt-get update
```

```
apt-get install -y docker.io
```

```
apt-get install -y kubelet kubeadm kubectl kubernetes-cni
```

# Multinode cluster setup

## Init Master node and Slave Nodes

- Run below command on master node.

```
kubeadm init --apiserver-advertise-address=<Master IP> --pod-network-cidr=10.244.0.0/16 --skip-preflight-checks
```

- Copy the join command that you will get after executing above init command.

Example: `kubeadm join --token 8ccfe0.7cd3dd5f91ef9796 <Master IP>:6443`

- Execute join command on all the slave nodes
- Start using cluster from master node

**\*\*Note\*\*** : We can do this from any node where kubectl installed

- Run below command to copy configuration to default directory

```
sudo cp /etc/kubernetes/admin.conf $HOME/
```

```
sudo chown $(id -u):$(id -g) $HOME/admin.conf
```

```
export KUBECONFIG=$HOME/admin.conf
```

- Setup Kubernetes Network. Run below command on master node

```
kubectl apply -f http://docs.projectcalico.org/v2.2/getting-started/kubernetes/installation/hosted/kubeadm/1.6/calico.yaml
```

# KIND

KIND is a tool for running local Kubernetes clusters using Docker container “nodes”.

It was primarily designed for testing Kubernetes itself, but may be used for local development or CI.

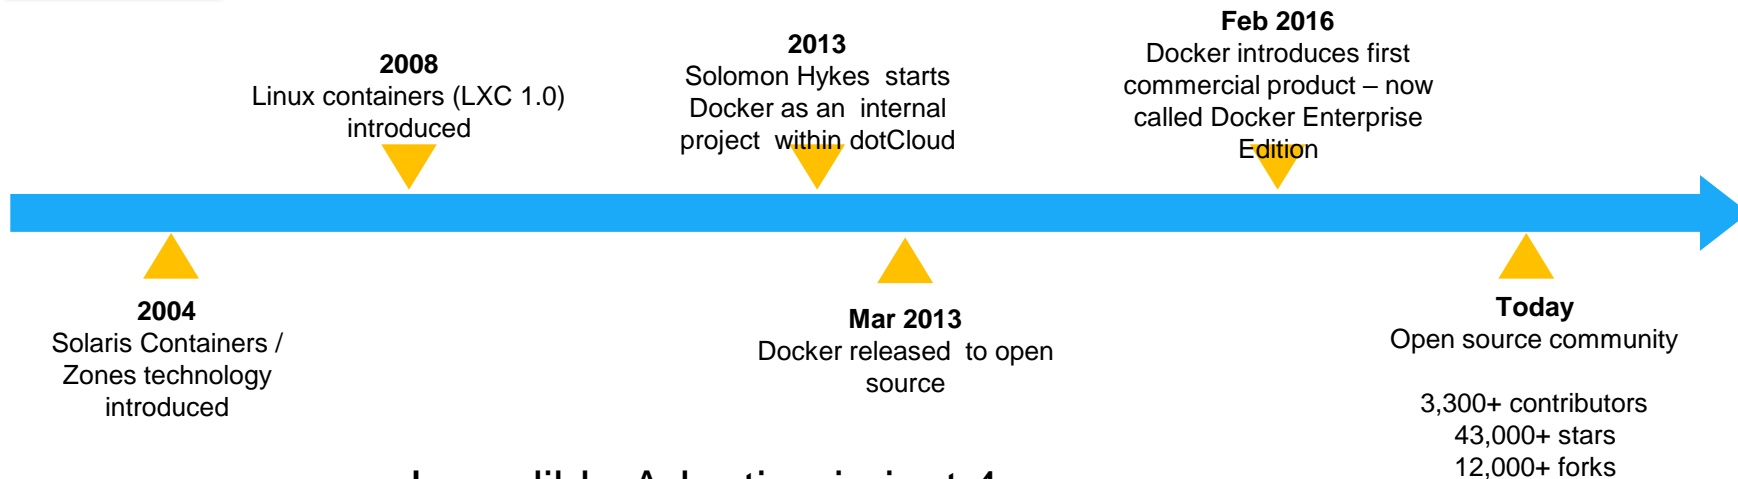
Kubernetes INside Docker

```
curl -Lo ./kind https://kind.sigs.k8s.io/dl/v0.8.1/kind-linux-amd64
```

```
chmod +x ./kind
```

```
sudo mv ./kind /usr/local/bin/kind
```

# History of Docker



Incredible Adoption in just 4 years



**14M**

Docker  
Hosts



**900K**

Docker  
apps



**77K%**

Growth in  
Docker job  
listings



**12B**


Image pulls  
Over 390K%  
Growth



**3300**

Project  
Contributors

# Continuous Integration and Delivery





# Tug-of-war



## Developers

- Freedom to create and deploy apps fast
- Define and package application needs



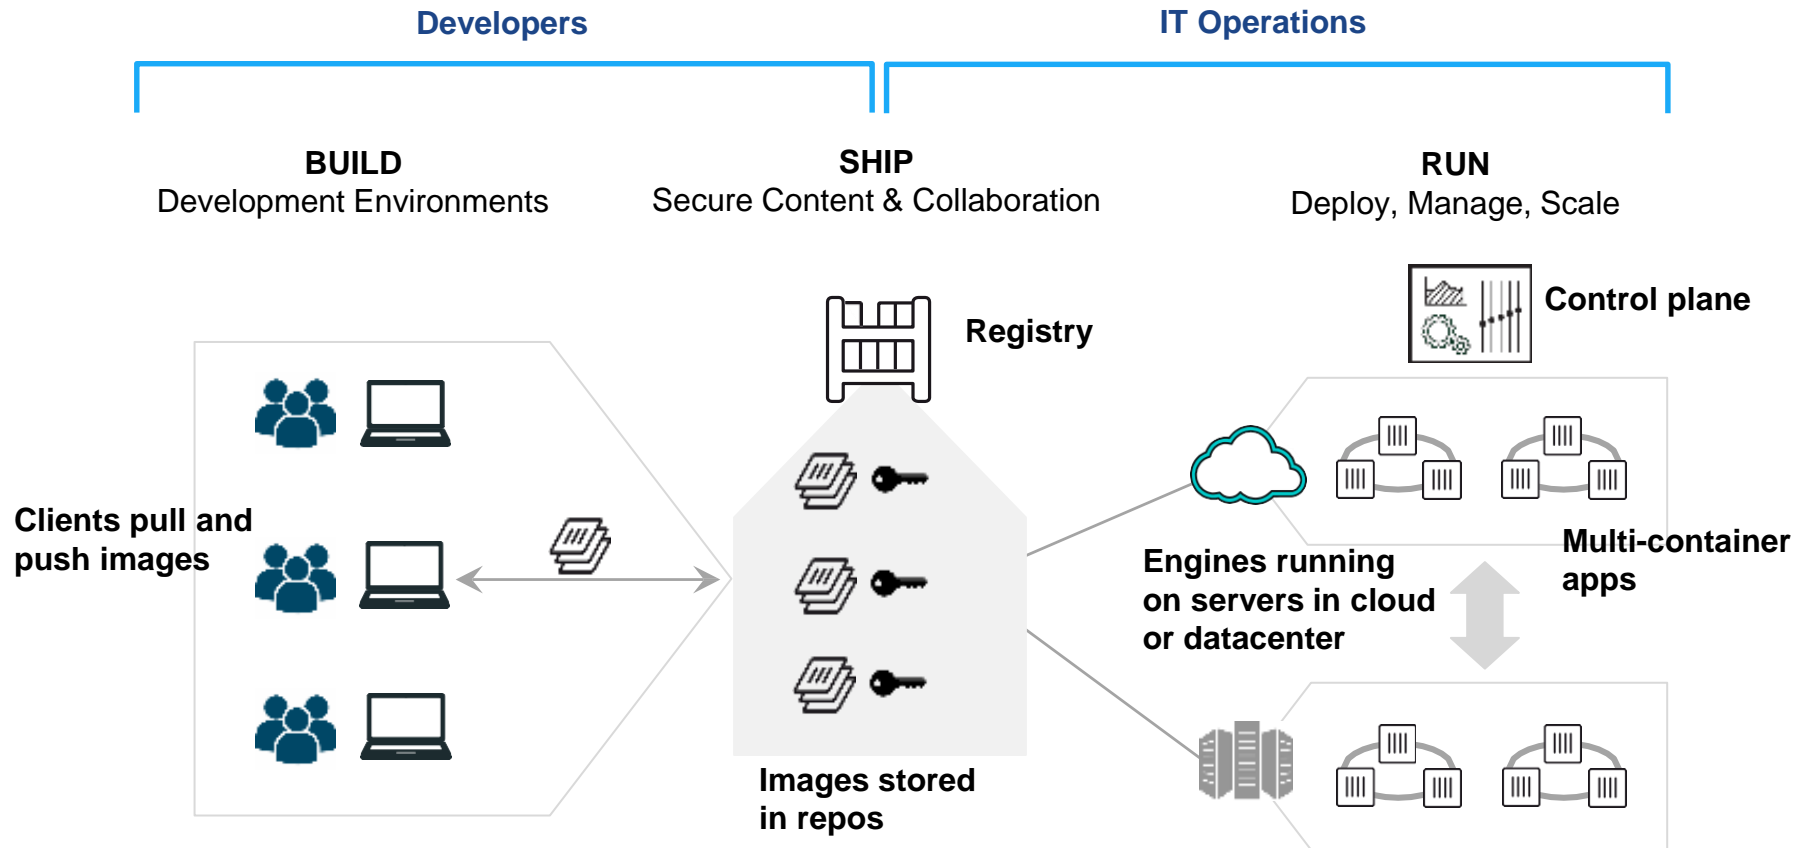
## IT Operations

- Quickly and flexibly respond to changing needs
- Standardize, secure, and manage

## The myth of bi-modal IT

	MICROSERVICES	TRADITIONAL APPS
Cloud or New Infrastructure	You are either here..	
Old Infrastructure		...or here

# Container as a service



# MobaXterm

Session settings



Warning: you have reached the maximum number of saved sessions for the personal edition of MobaXterm.  
You can start a new session but it will not be automatically saved.

Please support MobaXterm by subscribing to the Professional edition here: <https://mobaxterm.mobatek.net>

## Basic SSH settings

Remote host \*

☒ Specify username



Port

## Advanced SSH settings

### Terminal settings

### Network settings

### Bookmark settings

☒ X11-Forwarding

☒ Compression

Remote environment:

Execute command:

☐ Do not exit after command ends


SSH-browser type:

☐ Follow SSH path (experimental)

☒ Use private key

☐ Adapt locales on remote server

Execute macro at session start:



## Connect to instance [Info](#)

Connect to your instance i-01f9e5aaa5a303582 (k8s) using any of these options



EC2 Instance Connect

Session Manager

**SSH client**

Instance ID

 i-01f9e5aaa5a303582 (k8s)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is cognixia.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.  
 `chmod 400 cognixia.pem`
4. Connect to your instance using its Public DNS:  
 `ec2-15-206-163-49.ap-south-1.compute.amazonaws.com`

Example:

 `ssh -i "cognixia.pem" ubuntu@ec2-15-206-163-49.ap-south-1.compute.amazonaws.com`



**Note:** In most cases, the guessed user name is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI user name.